

Source: <https://www.kaggle.com/bmconrad/exploratory-analysis-stock-data>

Colab link: <https://colab.research.google.com/drive/1-fS-iu3PHhvHNO2kZhvrGeS-6sFyAx8T>

Another EDA version: <https://colab.research.google.com/drive/1gE7gf-b3ORjMMVogFePMfE8bzIzOGutV>

In [0]:

```
!nvidia-smi
```

Wed Nov 27 12:07:39 2019

```
+-----+
| NVIDIA-SMI 440.33.01      Driver Version: 418.67      CUDA Version: 10.1      |
+-----+-----+-----+-----+-----+-----+
| GPU   Name                Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan   Temp   Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|=====+=====+=====+=====+=====+=====+
|    0   Tesla T4              Off      | 00000000:00:04.0 Off |                    0 |
| N/A   70C    P8      12W / 70W | 0MiB / 15079MiB |      0%      Default |
+-----+-----+-----+-----+-----+-----+

+-----+-----+-----+-----+-----+-----+
| Processes:                                                       GPU Memory |
|  GPU       PID    Type    Process name                       Usage      |
|=====+=====+=====+=====+=====+=====+
|   No running processes found                                   |
+-----+-----+-----+-----+-----+-----+
```

In [0]:

```
# Importing libraries.
import pandas as pd
import numpy as np
import seaborn as sns #visualisation
import matplotlib.pyplot as plt #visualisation
%matplotlib inline
sns.set(color_codes=True)
from scipy.stats import norm
from subprocess import check_output
```

Load data

In [0]:

```
#install gdown
import os.path as path
if not path.exists('M3final.zip'):
    !pip install gdown
    !gdown https://drive.google.com/uc?id=1w4Sf7GXSzVs5Gcci9Dx2GqiT96A0OeeU
    !unzip M3final.zip
else :
    print('data is in places')
```

Requirement already satisfied: gdown in /usr/local/lib/python3.6/dist-packages (3.6.4)
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from gdown) (1.12.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.6/dist-packages (from gdown) (4.28.1)
Requirement already satisfied: requests in /usr/local/lib/python3.6/dist-packages (from gdown) (2.21.0)
Requirement already satisfied: urllib3<1.25,>=1.21.1 in /usr/local/lib/python3.6/dist-packages (from requests->gdown) (1.24.3)
Requirement already satisfied: chardet<3.1.0,>=3.0.2 in /usr/local/lib/python3.6/dist-packages (from requests->gdown) (3.0.4)
Requirement already satisfied: idna<2.9,>=2.5 in /usr/local/lib/python3.6/dist-packages (from requests->gdown) (2.8)
Requirement already satisfied: certifi<2017.4.17 in /usr/local/lib/python3.6/dist-packages

```
Requirement already satisfied: certifi==2017.4.17 in /usr/local/lib/python3.6/dist-packages (from requests->gdown) (2019.9.11)
Downloading...
From: https://drive.google.com/uc?id=1w4Sf7GXSzVs5Gcci9Dx2GqiT96A00eeU
To: /content/M3final.zip
31.8MB [00:00, 119MB/s]
Archive: M3final.zip
  inflating: fundamentals.csv
  inflating: __MACOSX/._fundamentals.csv
  inflating: prices-split-adjusted.csv
  inflating: __MACOSX/._prices-split-adjusted.csv
  inflating: prices.csv
  inflating: __MACOSX/._prices.csv
  inflating: securities.csv
  inflating: __MACOSX/._securities.csv
```

In [0]:

```
fund = pd.read_csv('fundamentals.csv')
price_split = pd.read_csv('prices-split-adjusted.csv')
prices = pd.read_csv('prices.csv')
comp_info = pd.read_csv('securities.csv')
#details of the data are described in stakeholders report
#in EDA part, we will use the price_split dataframe
```

In [0]:

```
#add new column named "change" which is the result of columns "open" and "close"
price_split['change'] = price_split['close'] - price_split['open']
price_split.head()
```

Out[0]:

	date	symbol	open	close	low	high	volume	change
0	2016-01-05	WLTW	123.430000	125.839996	122.309998	126.250000	2163600.0	2.409996
1	2016-01-06	WLTW	125.239998	119.980003	119.940002	125.540001	2386400.0	-5.259995
2	2016-01-07	WLTW	116.379997	114.949997	114.930000	119.739998	2489500.0	-1.430000
3	2016-01-08	WLTW	115.480003	116.620003	113.500000	117.440002	2006300.0	1.140000
4	2016-01-11	WLTW	117.010002	114.970001	114.089996	117.330002	1408600.0	-2.040001

Exploratory data analysis

First, we inspect the data and understand the dataframe before we do further analysis such as:

- What does some of our closing data look like?
- Who were the latest top 10 closers?
- What do their specs look like over time with respect to one another?
- What is their riskiness comparatively?

In [0]:

```
#inspect the datatype of the dataframe
price_split.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 851264 entries, 0 to 851263
Data columns (total 8 columns):
date           851264 non-null object
symbol         851264 non-null object
open           851264 non-null float64
close          851264 non-null float64
low            851264 non-null float64
high           851264 non-null float64
volume         851264 non-null float64
change         851264 non-null float64
```

```
dtypes: float64(6), object(2)
memory usage: 52.0+ MB
```

We notice that 'Date' column is displayed as an object and this can be problematic when we plot histograms/line chart so we convert it to a datetime datatype.

```
In [0]:
```

```
#turn 'Date' column to datetime
price_split["date"] = pd.to_datetime(price_split["date"])
```

```
In [0]:
```

```
#Here we use the date and close columns to have an overview over time
plt.figure(figsize=(20,8))
plt.plot('date','close',data=price_split)
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.xticks(rotation=45)
```


/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/converter.py:103: FutureWarning: Using an implicitly registered datetime converter for a matplotlib plotting method. The converter was registered by pandas on import. Future versions of pandas will require you to explicitly register matplotlib converters.

To register the converters:

```
>>> from pandas.plotting import register_matplotlib_converters
>>> register_matplotlib_converters()
warnings.warn(msg, FutureWarning)
```


```
Out[0]:
```


```
(array([733408., 733773., 734138., 734503., 734869., 735234., 735599.,
       735964., 736330.]), <a list of 9 Text xticklabel objects>)
```



```
In [0]:
```

```
#Next we look at stocks of the first 10 companies over time
symbols = price_split["symbol"].unique().tolist() #get the unique value from 'symbol' column and convert such series to list
for u in symbols[:10]: #we also use date and close columns
    dates = price_split[(price_split["symbol"] == u)]["date"]
    values = price_split[(price_split["symbol"] == u)]["close"]
    plt.plot(dates.tolist(), values.tolist())
plt.legend(symbols, loc='upper left')
plt.show()
```





Here we can see that...

In [0]:

```
#Mask by most recent year
latest_year = max(pd.unique(list(price_split["date"].apply(lambda x:x.year)))) #find the
most recent year which is 2016 only
latest_year_mask = [i==latest_year for i in price_split["date"].apply(lambda x:x.year)]
#create a filtered list of the stocks in 2016
df_masked1 = price_split[latest_year_mask] #apply such filter to price_split
df_masked1.head()
```

Out[0]:

	date	symbol	open	close	low	high	volume	change
0	2016-01-05	WLTW	123.430000	125.839996	122.309998	126.250000	2163600.0	2.409996
1	2016-01-06	WLTW	125.239998	119.980003	119.940002	125.540001	2386400.0	-5.259995
2	2016-01-07	WLTW	116.379997	114.949997	114.930000	119.739998	2489500.0	-1.430000
3	2016-01-08	WLTW	115.480003	116.620003	113.500000	117.440002	2006300.0	1.140000
4	2016-01-11	WLTW	117.010002	114.970001	114.089996	117.330002	1408600.0	-2.040001

In [0]:

```
#Mask by most top recent closers
#remind: symbols = price_split["symbol"].unique().tolist() (the same as above so we will
not run again)
maxPerSymbol={} #create a dictionary
for u in symbols:
    values = price_split[(price_split["symbol"] == u)]["close"].tolist() #get the close
price of companies and convert them to list
    maxPerSymbol[u] = max(values) #Get the names of companies from the dictionary having
the maximum close price
    #print(maxPerSymbol[u], u) #this step is optional to see the result

top10Closers = list(dict(sorted(maxPerSymbol.items()), #get the name of top 10 recent clos
ers
                        key=lambda v:v[1], #sort by key -> v[0] | sort by value->v[1]
                        reverse=True)[:10]).keys()) #sort in a descending order

top_10_mask = [i in top10Closers for i in df_masked1["symbol"].tolist()] #create filtere
d list of top 10 closes in 2016
df_masked2 = df_masked1[top_10_mask] #apply such filter to df_masked1 which is in 2016
df_masked2.head()
```

Out[0]:

	date	symbol	open	close	low	high	volume	change
725425	2016-01-04	AMZN	656.289978	636.989990	627.510010	657.719971	9314500.0	-19.299988
725441	2016-01-04	AZO	733.000000	735.479980	728.520020	742.229980	299100.0	2.479980
725452	2016-01-04	BIIB	300.339996	294.619995	291.399994	301.019989	2451100.0	-5.720001
725485	2016-01-04	CMG	468.700012	448.809998	447.500000	469.000000	2690300.0	-19.890014
725589	2016-01-04	GOOG	743.000000	741.840027	731.257996	744.059998	3272800.0	-1.159973


In [0]:

```
#Visualize top 10 most recent closers in 2016
symbols = df_masked2["symbol"].unique().tolist() #get the unique value from 'symbol' column from df_masked2 and convert such series to list
for u in symbols[:10]: #choose the first 10 top closers, we still use date and close columns
    dates = df_masked2[(price_split["symbol"] == u)]["date"]
    values = df_masked2[(price_split["symbol"] == u)]["close"]
    plt.plot(dates.tolist(), values.tolist())
#plt.legend(symbols, loc='upper left')
plt.legend(symbols, bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
plt.title(r'Top 10 Closing Stocks in 2016', fontsize=16)

plt.xlabel('Time (years)', fontsize=12)
plt.ylabel('Price ($/stock)', fontsize=12)
plt.show()
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

This is separate from the ipykernel package so we can avoid doing imports until
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:4: UserWarning: Boolean Series key will be reindexed to match DataFrame index.
after removing the cwd from sys.path.



The Priceline Group has the highest close price in 2016

In [0]:

```
#Visualize top 10 closers change. Now let's have a look at 'change' column
def build_density_facetwrap(somedf, colName, valName, wrapAmount, yourTitle):
    sns.set(style="white", rc={"axes.facecolor": (0, 0, 0, 0)})

    # Initialize the FacetGrid object
    pal = sns.cubehelix_palette(10, rot=-.25, light=.7)
    g = sns.FacetGrid(somedf, col=colName, hue=colName, col_wrap=wrapAmount, palette=pal)
    # Draw the densities in a few steps
    g.map(sns.kdeplot, valName, clip_on=False, shade=True, alpha=1, lw=1.5, bw=.2)
    g.map(sns.kdeplot, valName, clip_on=False, color="w", lw=0.5, bw=.2)
    g.map(plt.axhline, y=0, lw=2, clip_on=False)

    # Define and use a simple function to label the plot in axes coordinates
    def label(x, color, label):
        ax = plt.gca()
        ax.text(0, .4, label, fontweight="bold", color=color,
            ha="left", va="center", transform=ax.transAxes)
        plt.xlabel('Change ($/day)', fontsize=12)
        plt.ylabel('Frequency', fontsize=12)


    g.map(label, "change")
```

```
# Set the subplots to overlap
g.fig.subplots_adjust(hspace=0)

# Remove axes details that don't play well with overlap
g.set_titles("")
g.fig.suptitle(yourTitle, fontsize=32)
g.set(yticks=[])
g.despine(bottom=True, left=True)
g.fig.subplots_adjust(top=.8)
g.add_legend()

build_density_facetwrap(df_masked2,
                        "symbol",
                        "change",
                        5,
                        "Top 10 Closers Change Spread in 2016"
                        )
```

Top 10 Closers Change Spread in 2016



In [0]:

```
df_masked2["date"] = pd.to_datetime(df_masked2["date"]) #change the date to datetime
months = list(df_masked2["date"].apply(lambda x:x.month).unique()) #get the unique values
of months and turn them to list
df_masked2["month"] = df_masked2["date"].apply(lambda x: x.month) #create a month column
```

	date	symbol	open	...	volume	change	month
725425	2016-01-04	AMZN	656.289978	...	9314500.0	-19.299988	1
725441	2016-01-04	AZO	733.000000	...	299100.0	2.479980	1
725452	2016-01-04	BIIB	300.339996	...	2451100.0	-5.720001	1
725485	2016-01-04	CMG	468.700012	...	2690300.0	-19.890014	1
725589	2016-01-04	GOOG	743.000000	...	3272800.0	-1.159973	1
...
850964	2016-12-30	GOOGL	803.210022	...	1728300.0	-10.760010	12
851007	2016-12-30	ISRG	638.320007	...	267300.0	-4.150024	12
851079	2016-12-30	MTD	421.980011	...	124200.0	-3.420013	12
851115	2016-12-30	PCLN	1483.489990	...	405100.0	-17.429931	12
851147	2016-12-30	REGN	375.299988	...	612900.0	-8.209992	12

[2520 rows x 9 columns]

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

"""Entry point for launching an IPython kernel.

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

This is separate from the ipykernel package so we can avoid doing imports until

Selecting any 3 companies for visualizations on respective opening and closing stock prices.

In [0]:

```
comp_plot = comp_info.loc[(comp_info["Security"] == 'Yahoo Inc.') | (comp_info["Security"] == 'Microsoft Corp.') | (comp_info["Security"] == 'Facebook'), ["Ticker symbol"] ]["Ticker symbol"]
print(comp_plot)
```

```
181      FB
306     MSFT
Name: Ticker symbol, dtype: object
```



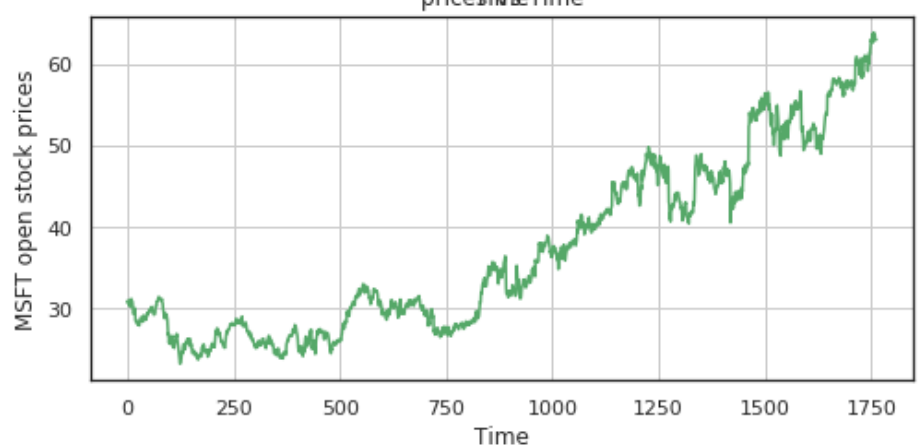

We notice that the differences is marginal , you might take open and close same but they are actually different after looking carefully.

In [0]:

```
def plotter(code):
    global closing_stock
    global opening_stock
    f, axs = plt.subplots(2,2,figsize=(8,8))
    plt.subplot(212)
    company = prices[prices['symbol']==code]
    company = company.open.values.astype('float32')
    company = company.reshape(-1, 1)
    opening_stock = company
    plt.grid(True)
    plt.xlabel('Time')
    plt.ylabel(code + " open stock prices")
    plt.title('prices Vs Time')
    plt.plot(company , 'g')

    plt.subplot(211)
    company_close = prices[prices['symbol']==code]
    company_close = company_close.close.values.astype('float32')
    company_close = company_close.reshape(-1, 1)
    closing_stock = company_close
    plt.xlabel('Time')
    plt.ylabel(code + " close stock prices")
    plt.title('prices Vs Time')
    plt.grid(True)
    plt.plot(company_close , 'b')
    plt.show()
for i in comp_plot:
    plotter(i)
```





Data Manufacturing

Feature Engineering

- The shape should be: |symbol| x |new features|
- Create a standard deviation for each stock
- Create a count for each positive day changes that occurred per stock

In [0]:

```
from __future__ import division

# Get each symbols standard deviation in time
ss = price_split.groupby(by=["symbol"])["change"].std()

# Count each symbol groups positive day change in time
## NOTE: it is a proportion to account for some days not appearing for some symbols in time
pcs = price_split.groupby('symbol').apply(lambda grp: grp[grp['change'] > 0]['change'].count() / grp['change'].size)

avgv = price_split.groupby(by=['symbol'])['volume'].mean()/10000000

newdf = pd.concat([ss, pcs, avgv], axis=1).reset_index() #create new dataframe
newdf.columns = ['symbol', 'std', 'prop_pos_day_change', "avg_volume"] #set the names for columns
newdf.head()
```

Out[0]:

	symbol	std	prop_pos_day_change	avg_volume
0	A	0.483919	0.520431	0.392759
1	AAL	0.583850	0.482974	0.935404
2	AAP	1.526761	0.513053	0.102203
3	AAPL	1.094628	0.507378	9.422578
4	ABBV	0.850244	0.542659	0.847079


Visualize our spread and proportion of GOOD change . i.e., Lower standard deviation and high positive change is the ideal stock choice.

In [0]:

```
for i in newdf['symbol'].tolist():
    x = newdf[newdf['symbol'] == i]['std']
    y = newdf[newdf['symbol'] == i]['prop_pos_day_change']
    plt.scatter(x,y)

legend = plt.legend(newdf['symbol'].tolist(),
                    bbox_to_anchor=(1.05, 1),
                    loc=2,
                    borderaxespad=0.,
                    ncol=10,
                    frameon = 1)
frame = legend.get_frame()
frame.set_color('white')

plt.title(r'Stock Change and Spread', fontsize=32)
plt.xlabel('Daily Positive Change (%)', fontsize=22)
plt.ylabel('Total Standard Deviation', fontsize=22)
plt.figure(figsize=(10,100))
plt.show()
```



Finished k= 5

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:6: FutureWarning: Method .as_matrix will be removed in a future version. Use .values instead.
```

Finished k= 6

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:6: FutureWarning: Method .as_matrix will be removed in a future version. Use .values instead.
```

Finished k= 7

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:6: FutureWarning: Method .as_matrix will be removed in a future version. Use .values instead.
```

Finished k= 8

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:6: FutureWarning: Method .as_matrix will be removed in a future version. Use .values instead.
```

Finished k= 9

Out[0]:

	symbol	std	prop_pos_day_change	k	values
0	A	0.483919	0.520431	3	0
1	AAL	0.583850	0.482974	3	0
2	AAP	1.526761	0.513053	3	1
3	AAPL	1.094628	0.507378	3	0
4	ABBV	0.850244	0.542659	3	0

Visualize Cluster Analysis

In [0]:

```
g = sns.FacetGrid(newdf, col="k", hue="values", col_wrap=4, palette='Set2')
g = g.map(plt.scatter, "std", "prop_pos_day_change")
g.set(xlabel="Closing Deviation")
g.set(ylabel="'Daily Positive Change (%)")
g.fig.suptitle("Stock Cluster Analysis", size=28)
g.fig.subplots_adjust(top=.8)
plt.subplots_adjust(hspace=1.2, wspace=0.4)
g.add_legend()
g._legend.set_title("Cluster")
```

